

December 1, 2022

ITM 885 Final Project

Classification Algorithms Comparison

Group 6: Talha Ali, Jorge Liakopulos, Di Ma

BROAD COLLEGE *of* BUSINESS

Table of Contents

1 Purpose of Project	Slide 3
2 Dataset	Slide 5
2.1 EDA	
2.2 Feature Engineering	
3 Model Building	
3.1 Random Forest	Slide 14
3.2 Logistic Regression	Slide 19
3.3 SVM	Slide 23
4 Conclusion	Slide 25

Purpose of Project

Classification Algorithms Comparison



Project Goal

Classification Algorithms Comparison

- Applying logistic regression, random forest, and SVM on the same dataset
- Comparing the performance
- Based on Titanic survival data

Dataset





	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S


Dataset

Passengers Survived the Titanic shipwreck

- What sorts of people were more likely to survive (0/1)
- 8 variables, including ticket class, sex, age, number of siblings/parents, fare, etc.
- Create models that predicts which passengers survived the Titanic shipwreck







White = Missing Values
Dark Blue = Non-Null



Feature Engineering

- Remove the column with more than 50% null values
- Fill the null values with mean
- Create new metrics by combining existing ones
 - For instance, create 'Family size' by combining parents and siblings variables

Random Forest



3.1 Random Forest

3.1.1 Dataset splitting

- Split the original dataset into train set and test set.

```
1 X = train.drop("Survived",axis=1)
2 Y = train["Survived"]
3 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=999)
```

3.1 Random Forest

3.1.2 Tuning the hyper-parameters of an estimator

- Use grid search to find out the best parameters
 - Comment out some parameters to use default values.

```
1 rfc = RandomForestClassifier(random_state=0)
2 rfc = rfc.fit(X_train,Y_train)
3 parameters = {'n_estimators': [5, 10, 20, 50],
4               #'max_features': ['auto'],
5               #'criterion': ['entropy', 'gini'],
6               #'max_depth': [2, 3, 5, 10],
7               #'min_samples_split': [],
8               #'min_samples_leaf': []
9            }
```

```
10 acc_scorer = make_scorer(accuracy_score)
11 grid = GridSearchCV(rfc, parameters, scoring=acc_scorer)
12 grid = grid.fit(X_train, Y_train)
13 rfc_best = grid.best_estimator_
14 rfc_best.fit(X_train, Y_train)
15 pred = rfc_best.predict(X_test)
16 acc_rf=accuracy_score(Y_test, pred)
17 acc_rf
```

0.8100558659217877

3.1 Random Forest

3.1.3 Performance

- 78% Accuracy of final result
- Top 18% ranking among all 13,000+ teams

AllSuccessfulErrors

Recent


Submission and DescriptionPublic Score

✓result_random_forest_csv


Complete · 3m ago

0.78229

2603madi1



0.7822911s



Your First Entry!


Welcome to the leaderboard! Your score represents your submission's accuracy. For example, a score of 0.7 in this competition indicates you predicted Titanic survival correctly for 70% of people.

What next? You've got a few options:

- 🧠 Learn skills that can improve your score in our [Intro to Machine Learning course](#) by Dan Becker.
- 🔍 Check out the [discussion forum](#) to find lots of tutorials and insights from other competitors.
- 🏆 Find a new challenge by entering one of our [open, active competitions](#) or searching our [public datasets](#).

3.1 Random Forest

3.1.3 Performance



3.1 Random Forest

3.1.3 Performance

	precision	recall	f1-score	support
0	0.85	0.90	0.87	549
1	0.83	0.74	0.78	342
accuracy			0.84	891
macro avg	0.84	0.82	0.83	891
weighted avg	0.84	0.84	0.84	891

Logistic Regression




3.2 Logistic Regression

```
X = train.drop("Survived",axis=1)
Y = train["Survived"]
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=999)
```

```
model = LogisticRegression(solver='liblinear',random_state=999).fit(X,Y)
model = model.fit(X_train, Y_train)
```

3.3 Logistic Regression



3.3 Logistic Regression

	precision	recall	f1-score	support
0	0.83	0.87	0.85	549
1	0.78	0.72	0.75	342
accuracy			0.81	891
macro avg	0.80	0.79	0.80	891
weighted avg	0.81	0.81	0.81	891

SVM



3.4 SVM

```
from sklearn.svm import SVC
classifier = SVC()
classifier.fit(X, Y)
score = classifier.score(X, Y)
print(score)
```

	precision	recall	f1-score	support
0	0.68	0.91	0.78	549
1	0.69	0.32	0.43	342
accuracy			0.68	891
macro avg	0.68	0.61	0.61	891
weighted avg	0.68	0.68	0.65	891